

ARC PRIZE 2025 RESEARCH-DRIVEN SOLVER DEVELOPMENT PROMPT

OVERVIEW AND MISSION

You are tasked with developing an AI systems to efficiently learn new skills and solve open-ended problems, rather than depend exclusively on systems trained with extensive datasets for the ARC Prize 2025 (Abstraction and Reasoning Corpus) competition on Kaggle. This is a research-driven development process where you will:

1. **Comprehensively review** all provided research documents and literature
2. **Synthesize** state-of-the-art approaches from the research
3. **Design** a novel or hybrid solution architecture based on your analysis
4. **Implement** the solution through iterative refinement
5. **Deliver** a production-ready Kaggle notebook (.py file)

Critical Principle: Your approach must be **derived from research**, not predetermined. You will analyze the literature, identify promising directions, and make informed architectural decisions based on empirical evidence and theoretical foundations from the research community.

****EXECUTION MODE**:** This is a single-pass, complete workflow. Execute all steps sequentially without pausing. For each step, produce the specified deliverables in full before proceeding to the next step.

****KAGGLE ENVIRONMENT**:**

- Your .py file is the solution code (what you submit to Kaggle)
- Input data path: `/kaggle/input/arc-prize-2025/arc-agi_test_challenges.json`
- Your code must write: `submission.json` in current directory
- Do NOT reference training/evaluation/solution files in the final .py code
- The .py code must be self-contained and work with only the test_challenges.json file

STEP 1: COMPREHENSIVE RESEARCH ANALYSIS AND PROBLEM UNDERSTANDING

Prompt:

Conduct a thorough analysis of the ARC Prize 2025 challenge by synthesizing all available resources in this directory.e.g everything in /home/legend/Documents/AGI/Kaggle, and sub folders like /home/legend/Documents/AGI/Kaggle/literature:

A. Document Review and Synthesis

1. ****Primary Problem Documents Analysis:****

- Read ALL provided files about the ARC Prize 2025
- Extract the exact problem specification:
 - * Task structure (JSON format, train/test pairs)
 - * Grid constraints (dimensions, color values)
 - * Submission format requirements (pass@2, attempt_1/attempt_2)
 - * Scoring methodology (exact matching)
- Identify competition constraints:
 - * Computational limits (12-hour runtime, no internet access)
 - * Hardware specifications (GPU availability, memory)
 - * Kaggle-specific requirements (file paths, deterministic execution)

2. ****Existing Research Literature Review:****

- Analyze the solutions so far and research documents provided (e.g.,
"/home/legend/Documents/AGI/Kaggle/literature/papers,
/home/legend/Documents/AGI/Kaggle/arc_prize_2025_submission and sun folders,
/home/legend/Documents/AGI/Kaggle/literature, any paper, code, approach, etc)
- Extract key methodologies mentioned: e.g
 - * Domain-Specific Language (DSL) approaches
 - * Program synthesis techniques
 - * LLM-augmented methods
 - * Object-centric neuro-symbolic approaches
- Identify what has worked historically:
 - * Performance benchmarks from literature and online searches
 - * Winning solutions from previous competitions
 - * Theoretical foundations cited

3. ****Online Research Supplementation:****

Use web_search to gather:

- Recent papers on ARC solving (2023-2025)
- Kaggle discussion forums for ARC Prize 2025
- Current leaderboard standings and publicly shared approaches

- GitHub repositories of successful ARC solvers
- Blog posts and technical write-ups from competition participants

B. Technical Problem Decomposition

Based on your research, answer:

1. ****What makes ARC tasks difficult?****

- Core cognitive requirements
- Why traditional ML/LLMs fail
- Specific reasoning challenges

2. ****What are the latest best proven solution paradigms?****

- Categorize approaches from literature (e.g symbolic, neural, hybrid etc)
- Identify success rates and limitations of each
- Note computational trade-offs

3. ****What are the key design decisions?****

- DSL expressiveness vs. search space size
- Search strategy (e.g breadth-first, beam search, Monte Carlo etc)
- Verification mechanisms
- Heuristic design
- Two-attempt strategy formulations

C. Comparative Analysis

Create a detailed comparison table:

- List 5-7 major approaches from literature
- For each: core principle, performance, computational cost, implementation complexity
- Identify gaps or opportunities for hybrid approaches

DELIVERABLE:

Research synthesis document (800-1000 words) including:

- Problem specification summary (200 words)
- Key findings from literature (300 words)
- Comparative analysis table (3-5 approaches)
- Preliminary architectural recommendation (200 words)

STEP 2: ARCHITECTURAL DESIGN AND APPROACH SELECTION

Prompt:

Based on your research synthesis, design a solver architecture:

A. Architecture Selection Rationale example: Use better one if you find a better one in your research

1. **Choose Your Core Approach:**

Justify your selection using research evidence:

- Will you use pure DSL-based program synthesis?
- Hybrid neuro-symbolic?
- LLM-augmented synthesis?
- Novel combination?

For your choice, provide:

- Supporting evidence from literature (cite specific papers/results)
- Why this approach suits ARC Prize 2025 constraints
- Expected performance range based on benchmarks
- Known limitations and mitigation strategies

2. **Component Architecture Design:**

Define the major system components:

- **Perception/Analysis Layer**: How will you analyze training pairs?
- **Reasoning/Synthesis Layer**: How will you generate candidate solutions?

- **Verification Layer**: How will you validate candidates?
- **Execution Layer**: How will you apply solutions to test inputs?

For each component:

- Technical approach (specific algorithms/methods)
- Research citations supporting the approach
- Expected computational complexity
- Failure modes and handling strategies

3. **Two-Attempt Strategy Design**:

Design how you'll generate diverse attempts:

- What defines "diversity" in your approach?
- How will attempt_1 and attempt_2 differ?
- Research evidence for ensemble/multi-strategy benefits

B. Implementation Specification

Create detailed specifications WITHOUT writing code:

1. **DSL Design (if applicable)**:

- List of primitive operations (justified by task analysis)
- Operation categories (geometric, color, structural, etc.)
- Parameter spaces for each operation
- Composability rules

2. **Search Strategy**:

- Algorithm choice (beam search, A*, MCTS, etc.)
- Search space pruning heuristics
- Stopping criteria
- Computational budgets per task

3. ****Heuristic Design:****

- Feature extraction from training pairs
- How features guide search
- Priority ordering mechanisms
- Adaptation strategies

4. ****Verification Mechanism:****

- Exact matching on training pairs
- Handling of ambiguous cases
- Confidence scoring (if applicable)

C. Theoretical Analysis

1. ****Completeness Analysis:****

- What percentage of ARC tasks is your approach theoretically capable of solving?
- What categories of tasks will it struggle with?

2. ****Computational Feasibility:****

- Estimated time per task
- Scalability to 100+ tasks
- Memory requirements

3. ****Expected Performance:****

- Predicted accuracy range based on literature benchmarks
- Comparison to current leaderboard (from online research)
- Conservative vs. optimistic estimates

DELIVERABLE:

Architecture design document (600-800 words):

- Approach selection rationale with citations (300 words)
- Component specifications (200 words)

- Theoretical performance analysis (150 words)

STEP 3: ITERATION 1 - INITIAL IMPLEMENTATION

Prompt:

Implement your first version of the solver based on your architectural design:

```
#### A. Implementation
```

Write a complete, production-ready Python file that:

1. Implements ALL components from your architecture
2. Loads data from Kaggle paths: `/kaggle/input/arc-prize-2025/arc-agi_test_challenges.json`

Guve kaggles notebook starter as provided by Kaggle {# This Python 3 environment comes with many helpful analytics libraries installed

It is defined by the kaggle/python Docker image:
<https://github.com/kaggle/docker-python>

For example, here's several helpful packages to load

```
import numpy as np # linear algebra
```

```
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

```
# Input data files are available in the read-only "../input/" directory
```

```
# For example, running this (by clicking run or pressing Shift+Enter) will  
list all files under the input directory
```

```
import os
```

```
for dirname, _, filenames in os.walk('/kaggle/input'):
```

```
    for filename in filenames:
```

```
        print(os.path.join(dirname, filename))
```

```
# You can write up to 20GB to the current directory (/kaggle/working/) that  
gets preserved as output when you create a version using "Save & Run All"
```

You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session}

3. Generates valid `submission.json` with correct format
4. Handles all edge cases gracefully
5. Runs deterministically (set seeds)
6. Includes clear code structure and comments

****Implementation Guidelines:****

- Follow your architecture specification exactly
- Implement complete functionality (no stubs or TODOs)
- Use only libraries available on Kaggle (numpy, scipy, no internet-dependent packages)
- Include error handling and fallback mechanisms
- Ensure computational efficiency where possible

B. Theoretical Correctness Analysis

Conduct focused self-review (300-400 words total):

- Logic verification: Walk through 2-3 example tasks mentally
- Compliance check: Verify all Kaggle requirements
- Critical assessment: 3 key strengths, 3 key weaknesses, performance estimate

DELIVERABLE:

- `arc_solver_v1.py`
- Self-assessment (300-400 words)

2. **Mental Task Walkthrough:**

Select 3 example task types (from your understanding of ARC):

- Simple pattern (e.g., color remapping)
- Geometric transformation (e.g., rotation + scaling)
- Complex compositional (e.g., multi-rule interaction)

For each, trace through your code mentally:

- Would it extract correct features?
- Would it generate appropriate candidates?
- Would verification work correctly?
- Would it produce valid output?

3. **Research Benchmark Comparison:**

- Compare your approach to methods in literature
- Identify where your implementation might underperform
- Identify potential advantages

C. Compliance and Feasibility Review

1. **Competition Requirements:**

- [] Correct input file path
- [] Correct submission.json format
- [] Two attempts per test input
- [] All tasks covered (no missing task_ids)
- [] Grids are valid (0-9 values, reasonable dimensions)
- [] No internet access required
- [] Deterministic execution

2. **Computational Complexity Analysis:**

- Estimate runtime per task (worst case, average case)
- Calculate total expected runtime for 100 tasks
- Identify computational bottlenecks
- Assess if it fits in 12-hour limit

3. **Code Quality:**

- Readability and maintainability
- Proper error handling

- Modularity and extensibility

D. Critical Self-Assessment

Write an honest critique:

- **Strengths**: What does this implementation do well?
- **Weaknesses**: Where is it likely to fail?
- **Theoretical Performance**: Estimate accuracy (e.g., "5-15%")
- **Comparison to SOTA**: How far from competitive solutions?
- **Key Improvements Needed**: Prioritized list

DELIVERABLE:

- `arc_solver_v1.py` (complete, runnable file)
- Self-assessment document (800-1200 words)
- Identified improvements for next iteration

STEP 4: ITERATION 2 - REFINEMENT AND ENHANCEMENT

Prompt:

Refine your solver based on V1 analysis and additional research:

A. Gap Analysis and Research

1. **Review V1 Weaknesses**:

- Identify the top 3-5 critical issues from V1 assessment
- For each issue, conduct targeted research:
 - * Use `web_search` for specific solutions (e.g., "ARC task color symmetry detection")
 - * Look for papers addressing these specific challenges
 - * Search Kaggle discussions for practical solutions

2. ****Enhancement Opportunities:****

- Research advanced techniques not in V1:
 - * More sophisticated search strategies
 - * Better heuristics
 - * Additional DSL operations
 - * Improved verification methods
- Find specific algorithmic improvements from recent work

B. Implementation V2

Create an improved version:

1. ****Address Critical Issues:****

- Fix logical errors from V1
- Improve components that were weak
- Add missing functionality

2. ****Add Enhancements:****

- Implement research-backed improvements
- Expand DSL (if applicable) with new operations
- Improve search efficiency
- Better feature extraction

3. ****Maintain Strengths:****

- Keep working components from V1
- Preserve successful design decisions
- Ensure compliance still met

C. Advanced Evaluation

1. ****Cross-Validation Against Research:****

- Compare V2 approach to 3-5 recent papers
 - Identify which techniques from literature you've incorporated
 - Note which proven techniques you're still missing
2. **Theoretical Task Coverage:**
- Categorize ARC task types (from problem understanding)
 - For each category, assess: Can V2 solve it?
 - Estimate coverage: "V2 should handle X% of Y-type tasks"
3. **Computational Profile:**
- Re-analyze runtime complexity
 - Compare to V1: Is it faster/slower?
 - Is it still within 12-hour budget?
4. **Online Benchmarking:**
- Use web_search to find:
 - * Current Kaggle leaderboard scores
 - * Discussion of what scores are competitive
 - * Public solution approaches and their performance
 - Compare V2's expected performance to these benchmarks

D. Critical Assessment V2

1. **Progress Evaluation:**
- Quantify improvements over V1
 - Justify why changes should help
 - Estimate new performance range
2. **Remaining Gaps:**
- What still doesn't work?
 - What tasks will V2 still fail on?

- Why?

3. ****Path to V3:****

- Identify THE most impactful improvement for V3
- Research backing for this improvement
- Implementation strategy

DELIVERABLE:

- `arc_solver_v2.py`
- V1→V2 improvement summary (400-500 words)
- Performance estimate with justifications

STEP 5: ITERATION 3 - OPTIMIZATION AND ROBUSTNESS

Prompt:

Create the most robust and competitive version before finalization:

A. State-of-the-Art Alignment Research

1. ****Deep Dive into Top Approaches:****

- Use web_search to find:
 - * Recent competition winners' approaches (2024-2025)
 - * Highest-performing published methods
 - * Novel techniques from latest papers
- Extract specific implementation details
- Identify "secret sauce" components

2. ****Competitive Gap Analysis:****

- Current leaderboard scores (from online research e.g <https://www.kaggle.com/competitions/arc-prize-2025/leaderboard>)
- Your V2 estimated performance

- Gap size and components causing the gap
- Feasibility of closing gap with V3

B. Implementation V3

Final pre-production version:

1. ****Incorporate SOTA Techniques:****
 - Add most promising technique from research
 - Implement sophisticated optimizations
 - Enhance robustness and edge case handling
2. ****Optimization Pass:****
 - Profile computational bottlenecks (theoretically)
 - Implement caching where beneficial
 - Optimize critical paths
 - Balance quality vs. speed
3. ****Robustness Hardening:****
 - Comprehensive error handling
 - Graceful degradation for hard tasks
 - Ensure NO crashes regardless of input
 - Validate all outputs before submission

C. Comprehensive Evaluation

1. ****Multi-Dimensional Analysis:****
 - **Correctness:****
 - Mental walkthrough of 5+ diverse task types
 - Verify logic for each component

- Check for subtle bugs or edge cases

****Compliance:****

- Re-verify all Kaggle requirements
- Check submission format rigorously
- Confirm deterministic behavior

****Performance:****

- Estimate accuracy on:
 - * Simple tasks (pattern matching)
 - * Medium tasks (2-3 operation compositions)
 - * Complex tasks (4+ operations, interactions)
- Overall accuracy estimate

****Efficiency:****

- Task time distribution (fast/slow tasks)
- Total expected runtime
- Memory usage assessment

2. ****Research-Based Validation:****

- Compare V3 to each major approach from Step 1 literature review
- For each: What techniques did you adopt? What did you skip? Why?
- Identify your solver's unique contributions or combinations

3. ****Online Competitive Analysis:****

- Latest Kaggle leaderboard standings
- Discussion forum insights
- Comparison: Where does V3 likely rank?
- What would push it higher?

D. Failure Mode Analysis

1. ****Identify Task Categories V3 Will Fail:****

- Specific examples from ARC task taxonomy
- Why will it fail? (insufficient DSL, search timeout, etc.)
- Could these be addressed? At what cost?

2. ****Known Limitations:****

- Computational constraints causing shortcuts
- Theoretical limitations of approach
- Trade-offs made and their implications

3. ****Risk Assessment:****

- Probability of submission errors
- Probability of timeout
- Expected score variance (best/worst case)

E. V3 Critical Assessment

Final honest evaluation:

- ****Architecture Quality****: Is the design sound?
- ****Implementation Quality****: Is the code production-ready?
- ****Expected Performance****: Realistic accuracy estimate with reasoning
- ****Competitive Standing****: Where in leaderboard range?
- ****Confidence Level****: High/Medium/Low and why
- ****Key Strengths****: Top 3 advantages
- ****Key Weaknesses****: Top 3 limitations

DELIVERABLE:

- `arc_solver_v2.py`
- V1→V2 improvement summary (400-500 words)
- Performance estimate with justification

STEP 6: FINAL VERSION - PRODUCTION IMPLEMENTATION

Prompt:

Create the final, production-ready solver for Kaggle submission:

A. Final Implementation Decision

Based on V3 assessment:

1. ****If V3 is strong****: Polish and finalize it
2. ****If critical issue found****: Make targeted fix, creating V3.1
3. ****If major rework needed****: Explain what and why, then implement

B. Production-Ready Code

Create `arc_solver_final.py`:

1. ****Code Quality****
 - Clean, well-commented code
 - Proper structure and organization
 - Professional naming conventions
 - Comprehensive docstrings
2. ****Robustness****
 - Handles ALL edge cases
 - Never crashes or hangs
 - Always produces valid submission
 - Graceful degradation for hard tasks
3. ****Kaggle Optimization****
 - Efficient imports (only necessary libraries)

- Proper file paths for Kaggle environment
- Deterministic with fixed random seeds
- Appropriate logging/progress indicators

4. **Submission Guarantee:**

- Always generates valid submission.json
- Correct format for every task
- Two attempts per test input
- No missing or malformed entries

C. Final Validation

1. **Mental Execution Test:**

- Trace through the complete pipeline
- Verify data flow from input to submission
- Check all branches and conditions
- Confirm error handling paths work

2. **Requirements Checklist:**

- [] Correct Kaggle input path
- [] Valid submission.json output
- [] Pass@2 format (two attempts)
- [] No internet dependencies
- [] Deterministic execution
- [] Runs within time limit (estimated)
- [] Handles all 100+ tasks
- [] Values in [0-9] range
- [] Grid dimensions valid (1x1 to 30x30)

3. **Final Complexity Analysis:**

- Best case runtime

- Average case runtime
- Worst case runtime
- Confidence in 12-hour completion

D. Documentation Package

Create concise documentation:

1. ****Approach Evolution**** (600 words):
 - Research foundation and key decisions (250 words)
 - Version progression V1→V2→V3→Final (250 words)
 - Implementation highlights (100 words)
2. ****Final Approach Explanation**** (600 words):
 - Architecture and core components (250 words)
 - Algorithmic strategy and two-attempt design (200 words)
 - Expected performance and limitations (150 words)

DELIVERABLE:

- `arc_solver_final.py`
- `approach_documentation.md` (both sections combined, ~1200 words total)

STEP 7: FINAL RESEARCH REFLECTION AND SUBMISSION PACKAGE

Prompt:

STEP 7: FINAL PACKAGE

Organize all deliverables with brief summary:

- Final performance prediction (confidence intervals)
- List of all 4 Python files produced
- 200-word reflection on research-to-implementation journey

DELIVERABLE:

- Summary document (3000 words)